

EffDim CPU and GPU benchmark results

Executive summary

EffDim was benchmarked on seeded standard-normal arrays with 768 features at 10,000, 50,000, and 100,000 samples.

The original CPU implementation took 2h 50m and peaked at 78 GiB for 100,000 samples. Most of this cost came from exact CPU nearest-neighbor work and GMST’s independent dense $N \times N$ distance matrix. After excluding GMST and moving k-NN to cuVS, the remaining 15 metrics completed in 12.9 seconds with exact neighbors.

For this all-points query workload, exact GPU k-NN was faster than either approximate backend. CAGRA and IVF-Flat remain candidates for much larger datasets, but their indexing and traversal overhead did not pay off through 100,000 samples.

On three real UniverseTBD embedding datasets, CAGRA achieved 99.04-99.99% recall. Rust regular PCA was $4.52\times$ faster than Python on the $100k \times 1024$ Legacy Survey matrix, while Python streaming covariance was $5.37\times$ faster than the current Rust streaming implementation. The fastest end-to-end non-GMST configuration on that dataset was Python streaming PCA plus shared CAGRA at 23.49 seconds.

Benchmark environment

- Host: angus-MS-7C56
- CPU: AMD Ryzen 7 5700X3D, 8 cores / 16 logical CPUs
- RAM: 125 GiB
- GPU: NVIDIA RTX PRO 6000 Blackwell Max-Q, 98 GB VRAM
- GPU compute capability: 12.0
- Python: 3.12.3
- Rust: 1.97.1
- Branch: rust-migration
- Base source commit: d60ef97
- Input: seeded standard-normal arrays plus selected UniverseTBD embeddings
- Execution: sequential; only one benchmark process ran at a time
- Repetitions: one measured pass per configuration, with no warmup

The benchmark bindings and scripts were working-tree additions on top of the listed source commit.

1. Original CPU full suite

This benchmark ran all 16 metrics, including exact CPU k-NN and Euclidean GMST. The Rust extension was compiled in release mode.

Samples	Compute time	Peak RSS	PCA dimensions at 95%
10,000	100.7 s (1m 41s)	1.12 GiB	703
50,000	2,529.5 s (42m 10s)	20.47 GiB	720
100,000	10,179.3 s (2h 49m 39s)	77.96 GiB	723

The three sequential runs took approximately 3h 33m. Runtime increased by about $101\times$ when sample count increased by $10\times$, closely matching $O(N^2)$ scaling.

2. Spectral/PCA-only scaling

The spectral path omits all nearest-neighbor and geometry estimators.

Samples	PCA/spectral time	Peak RSS	PCA dimensions at 95%
10,000	0.62 s	443 MiB	703
50,000	1.78 s	2.03 GiB	720

Samples	PCA/spectral time	Peak RSS	PCA dimensions at 95%
100,000	3.60 s	4.11 GiB	723

One-million-row streaming PCA follow-up

A direct $1,000,000 \times 768$ float64 run used a seeded standard-normal matrix, 4,096-row chunks, 16 threads, and three repetitions. The 5.72 GiB input was memory-mapped and reused by isolated workers.

Implementation	Best time	PCA dimensions at 95%	Peak RSS
Python NumPy/OpenBLAS	3.437 s	728	5.82 GiB
Rust ndarray/OpenBLAS	4.154 s	728	5.80 GiB
Rust faer GEMM	6.354 s	728	5.81 GiB

This validates one-million-row streaming PCA on the host. Rust OpenBLAS was $1.21\times$ slower than Python and $1.53\times$ faster than faer. Peak RSS is dominated by resident pages from the memory-mapped input, rather than PCA workspace alone.

3. Isolated GPU nearest-neighbor experiments

NVIDIA cuVS 26.6 with CUDA 13 was tested using exact brute force, CAGRA, and IVF-Flat. Recall was measured against exact GPU neighbors on 1,000 queries.

Exact cuVS k-NN

Samples	All-points k-NN time	Recall@10
10,000	0.0148 s	100%
50,000	0.324 s	100%
100,000	1.40 s	100%

ANN trade-offs at 100,000 samples

Backend and setting	Recall@10	All-points search	Peak VRAM
CAGRA, itopk=128	40.22%	0.90 s	1.72 GiB
CAGRA, itopk=1024	91.34%	7.99 s	6.30 GiB
CAGRA, itopk=2048	98.19%	15.88 s	10.80 GiB
IVF-Flat, nprobe=128/316	57.58%	4.67 s	2.45 GiB
IVF-Flat, nprobe=256/316	91.00%	9.51 s	2.23 GiB
IVF-Flat, all 316 lists	100%	6.86 s	2.23 GiB
Exact cuVS brute force	100%	1.40 s	about 1.4 GiB

Random Gaussian vectors in 768 dimensions are difficult for ANN because distances are highly concentrated. Recall and speed on real embedding data may differ materially.

The available FAISS CUDA wheel installed successfully but failed at runtime with CUDA error 209: no kernel image is available for execution on the device. That wheel did not contain kernels for the RTX workstation GPU’s sm_120 architecture. cuVS CUDA 13 worked correctly.




Figure 1: Original full-suite RSS




Figure 2: PCA-only RSS

4. GPU-backed 15-metric suite, GMST excluded

This benchmark replaced Rust’s exact CPU k-NN with each cuVS backend. It ran the eight spectral metrics and seven k-NN-based geometry metrics. GMST was excluded because it separately constructs dense $N \times N$ distance matrices and would hide the effect of changing the k-NN backend.

CAGRA was tuned for approximately 98% recall. IVF-Flat searched 80% of its coarse lists and achieved approximately 90% recall.

Samples	Backend	Total time	k-NN time	Recall@10	Peak RSS	Peak VRAM
10,000	Exact cuVS	1.71 s	0.015 s	100%	0.89 GiB	1.10 GiB
10,000	CAGRA	2.54 s	0.107 s	99.34%	1.37 GiB	1.72 GiB
10,000	IVF-Flat	1.92 s	0.059 s	89.44%	1.11 GiB	1.30 GiB
50,000	Exact cuVS	6.16 s	0.318 s	100%	2.78 GiB	1.22 GiB
50,000	CAGRA	9.10 s	2.394 s	97.98%	3.18 GiB	5.02 GiB
50,000	IVF-Flat	7.30 s	1.218 s	89.71%	2.86 GiB	1.68 GiB
100,000	Exact cuVS	12.91 s	1.379 s	100%	5.26 GiB	1.37 GiB
100,000	CAGRA	28.48 s	15.910 s	98.26%	5.44 GiB	10.80 GiB
100,000	IVF-Flat	20.97 s	9.206 s	89.91%	5.29 GiB	2.22 GiB




Figure 3: GPU-backed suite memory

At 100,000 samples, the CAGRA estimates remained close to exact despite 98.26% neighbor recall. For example, MLE dimensionality differed by 0.46% and Two-NN differed by 0.21%. IVF-Flat’s approximately 90% recall produced MLE and Two-NN differences below 1% on this synthetic dataset.

Conclusions

1. Exact cuVS is the best tested k-NN backend through 100,000 samples. Its batched all-points workload efficiently saturates the GPU, while ANN index traversal adds more overhead than it saves.

2. CAGRA provides a useful high-recall fallback for datasets where exact GPU search eventually becomes too expensive. At 100,000 samples, however, its 98% recall configuration was more than $11\times$ slower than exact GPU k-NN.
3. IVF-Flat used less VRAM than high-recall CAGRA but its tested 80%-probe configuration delivered only about 90% recall and was still slower than exact search.
4. Replacing k-NN alone is insufficient for the original 16-metric suite. GMST’s dense pairwise matrix is the remaining quadratic memory bottleneck and needs a sparse, batched, sampled, or separately reported implementation.
5. A one-million-sample full-suite run is not practical until GMST is redesigned. Streaming PCA is demonstrated at one million samples; the GPU-backed non-GMST suite remains the next full-pipeline scaling candidate.

5. Python main versus Rust on UniverseTBD embeddings

Scope

This comparison uses the current pure-Python implementation from main at commit d1e7af9 and the Rust migration based on commit d60ef97. Both implementations consume the **same CAGRA distances and neighbor indices**. GMST is excluded. Streaming PCA means chunked Chan covariance accumulation followed by covariance eigendecomposition, using float64 accumulation and a 4,096-row chunk size in both languages.

All runs were sequential and used one measured pass. Timings therefore include normal run-to-run noise; sub-millisecond scalar metric timings should be treated as directional rather than as stable microbenchmarks.

Real datasets

Dataset	Shape	CAGRA build	CAGRA search	Recall@10
JWST	$1,496 \times 1,024$	0.836 s	0.033 s	99.55%
DESI	$20,465 \times 768$	0.781 s	0.576 s	99.04%
Legacy Survey	$100,000 \times 1,024$	0.902 s	15.312 s	99.99%

The Legacy Survey input is capped at 100,000 rows. Only the selected embedding columns were streamed from Hugging Face; the full 108 GB repository was not downloaded.

PCA comparison

Dataset	Python regular	Rust regular	Advantage	Python streaming	Rust streaming	Advantage
JWST	1.817 s	0.175 s	$10.37\times$ Rust	0.061 s	0.220 s	$3.59\times$ Python
DESI	0.436 s	1.057 s	$2.42\times$ Python	0.106 s	0.593 s	$5.61\times$ Python
Legacy Survey	23.163 s	5.119 s	$4.52\times$ Rust	0.932 s	5.003 s	$5.37\times$ Python

Dataset	Python PCA-95	Rust PCA-95	Python peak RSS	Rust peak RSS
JWST	145	145	0.27 GiB	0.09 GiB
DESI	64	64	1.56 GiB	0.82 GiB
Legacy Survey	110	110	9.54 GiB	5.15 GiB

PCA result: Rust regular SVD wins decisively on the small JWST and large Legacy Survey matrices, but NumPy wins on the medium DESI shape. Python streaming covariance is currently




Figure 4: Python and Rust PCA comparison

faster on all three datasets in this table. That run predates the OpenBLAS follow-up: linking `ndarray::dot` to OpenBLAS later reduced the $100k \times 1,024$ Rust streaming result from 5.00 s to 1.30 s, versus 0.74 s for Python.

Python main switches to randomized SVD when both matrix dimensions are at least 1,000, while Rust regular PCA remains exact. Streaming covariance provides an exact common reference and avoids that algorithm difference.

PCA numerical agreement

Dataset	Python regular vs streaming	Rust regular vs streaming	Python vs Rust streaming
JWST	2.50e-14	4.93e-05	3.42e-02
DESI	1.48e-10	2.89e-10	4.38e-10
Legacy Survey	3.25e-02	1.89e-05	1.89e-04

Values are the maximum relative difference across the eight spectral dimensionality outputs. The larger Python regular-path differences on the 1,024-feature datasets come from its randomized-SVD branch and omission of the final component.

End-to-end non-GMST pipeline

These totals add shared CAGRA build/search, PCA, scalar spectral metrics, and the seven geometry estimators. Rust uses its bundled geometry path so data conversion and neighbor arrays are shared once.

Dataset	Python regular total	Rust regular total	Python streaming total	Rust streaming total
jwst_dinov3_vitl16	2.824 s	1.175 s	1.068 s	1.220 s
desi_dinov3_small_vitl16	2.765 s	3.784 s	2.435 s	3.321 s
legacysurvey_dinov3_vitl16	45.724 s	30.411 s	23.493 s	30.294 s

Per-estimator comparison

jwst_dinov3_vitl16

Estimator	Python	Rust	Speed advantage	Relative value difference
PCA explained variance (95%)	44.8 μ s	7.7 μ s	5.82 $\times$ Rust	0.00e+00
Participation ratio	14.5 μ s	3.2 μ s	4.52 $\times$ Rust	3.48e-15
Shannon entropy	33.2 μ s	10.8 μ s	3.09 $\times$ Rust	5.31e-15
Rényi $\alpha=2$	7.7 μ s	17.5 μ s	2.27 $\times$ Python	6.05e-16
Rényi $\alpha=3$	34.4 μ s	12.2 μ s	2.82 $\times$ Rust	3.62e-15
Rényi $\alpha=4$	23.7 μ s	12.2 μ s	1.94 $\times$ Rust	5.67e-15
Rényi $\alpha=5$	23.5 μ s	12.1 μ s	1.94 $\times$ Rust	4.73e-15
Geometric mean	48.2 μ s	8.2 μ s	5.87 $\times$ Rust	3.43e-02
MLE	202.7 μ s	445.3 μ s	2.20 $\times$ Python	8.21e-08
Two-NN	121.9 μ s	248.6 μ s	2.04 $\times$ Python	3.14e-08
DANCo	91.02 ms	107.33 ms	1.18 $\times$ Python	9.27e-10
MiND-MLi	74.8 μ s	205.3 μ s	2.74 $\times$ Python	2.62e-08
MiND-MLk	198.6 μ s	300.4 μ s	1.51 $\times$ Python	2.42e-07
ESS	45.36 ms	24.63 ms	1.84 $\times$ Rust	1.96e-08
TLE	138.8 μ s	303.0 μ s	2.18 $\times$ Python	1.46e-07

desi_dinov3_small_vitl16

Estimator	Python	Rust	Speed advantage	Relative value difference
PCA explained variance (95%)	43.2 μ s	8.9 μ s	4.83 $\times$ Rust	0.00e+00
Participation ratio	14.4 μ s	3.5 μ s	4.10 $\times$ Rust	2.87e-15
Shannon entropy	27.9 μ s	12.2 μ s	2.28 $\times$ Rust	3.53e-15
Rényi $\alpha=2$	7.1 μ s	15.1 μ s	2.12 $\times$ Python	6.56e-15
Rényi $\alpha=3$	29.6 μ s	11.4 μ s	2.60 $\times$ Rust	3.74e-15
Rényi $\alpha=4$	19.1 μ s	11.3 μ s	1.69 $\times$ Rust	3.39e-15
Rényi $\alpha=5$	19.0 μ s	11.3 μ s	1.69 $\times$ Rust	2.83e-15
Geometric mean	37.6 μ s	9.0 μ s	4.18 $\times$ Rust	4.83e-15
MLE	1.36 ms	35.37 ms	26.04 $\times$ Python	3.20e-10
Two-NN	483.2 μ s	34.45 ms	71.29 $\times$ Python	4.27e-08
DANCo	590.44 ms	1.114 s	1.89 $\times$ Python	4.41e-08
MiND-MLi	157.0 μ s	33.26 ms	211.84 $\times$ Python	7.91e-08
MiND-MLk	1.23 ms	33.21 ms	27.07 $\times$ Python	4.04e-07
ESS	366.98 ms	286.98 ms	1.28 $\times$ Rust	1.09e-07
TLE	1.02 ms	32.83 ms	32.24 $\times$ Python	3.20e-10

legacysurvey_dinov3_vitl16

Estimator	Python	Rust	Speed advantage	Relative value difference
PCA explained variance (95%)	42.4 μ s	9.2 μ s	4.61 $\times$ Rust	0.00e+00
Participation ratio	13.3 μ s	3.6 μ s	3.74 $\times$ Rust	3.65e-15
Shannon entropy	28.9 μ s	14.7 μ s	1.96 $\times$ Rust	6.70e-15
Rényi $\alpha=2$	6.9 μ s	15.4 μ s	2.22 $\times$ Python	2.11e-15
Rényi $\alpha=3$	34.8 μ s	12.2 μ s	2.84 $\times$ Rust	1.91e-15
Rényi $\alpha=4$	24.1 μ s	12.6 μ s	1.92 $\times$ Rust	1.31e-15
Rényi $\alpha=5$	23.2 μ s	12.1 μ s	1.91 $\times$ Rust	1.26e-15

Estimator	Python	Rust	Speed advantage	Relative value difference
Geometric mean	39.2 μ s	8.9 μ s	4.42 $\times$ Rust	3.27e-02
MLE	9.40 ms	236.33 ms	25.15 $\times$ Python	4.38e-08
Two-NN	1.95 ms	220.57 ms	113.39 $\times$ Python	7.01e-08
DANCo	3.859 s	7.344 s	1.90 $\times$ Python	2.09e-09
MiND-MLi	602.8 μ s	218.40 ms	362.31 $\times$ Python	8.25e-08
MiND-MLk	6.00 ms	228.30 ms	38.07 $\times$ Python	3.31e-07
ESS	2.402 s	1.929 s	1.25 $\times$ Rust	9.85e-08
TLE	4.82 ms	230.40 ms	47.81 $\times$ Python	4.38e-08

Interpretation

1. **CAGRA is accurate on real embeddings.** Recall@10 ranges from 99.04% to 99.99%, substantially better than on isotropic Gaussian vectors at comparable settings.
2. **Rust regular PCA is shape-dependent but strong at scale.** It is about 4.5 $\times$ faster than Python regular PCA on the $100k \times 1024$ Legacy Survey matrix, while NumPy is faster on the $20k \times 768$ DESI matrix.
3. **Python streaming covariance remains fastest.** In the original run, NumPy was roughly 5 $\times$ faster on the two larger datasets. The subsequent OpenBLAS-backed Rust path narrowed the $100k \times 1,024$ result to 1.30 s versus 0.74 s for Python, a remaining 1.77 $\times$ gap.
4. **Estimator values agree closely.** Geometry outputs use identical CAGRA neighbors and agree to numerical precision. Differences in regular spectral results primarily reflect Python’s randomized-SVD branch, not a language-level discrepancy.
5. **Cheap per-method Rust calls pay conversion overhead.** Several simple geometry estimators look slower through individual PyO3 calls because each call converts the full input to float32 and copies precomputed distances. The bundled Rust geometry path shares those conversions and is the relevant end-to-end implementation.