
Where Does Platonic Convergence Live? Pointwise Curvature of a Galaxy Embedding Manifold and What a Linear Map Can See

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The Platonic Universe study measured the convergence of foundation-model rep-
2 presentations on galaxy images as one number per pair of embeddings. We ask
3 whether that convergence is concentrated where the embedding manifold is flat.
4 A pointwise curvature field, obtained by differentiating a trained decoder and vali-
5 dated against analytic truth in that regime, finds mean-curvature magnitude corre-
6 lating negatively with cross-survey neighbour agreement at every latent dimension,
7 with local density explaining 50 to 80 percent of it. A second-order expansion of
8 a linear map restricted to the manifold explains why: a readout sees the second
9 fundamental form contracted with its normal component, $\langle w_N, \Pi \rangle$, and an align-
10 ment sees the relative second fundamental form of the two embeddings, neither
11 visible to mean-curvature magnitude at large codimension. On a unit-normalized
12 embedding that Hessian splits into a sphere term, proportional to the readout’s own
13 value, and a shape term. On a synthetic surface at production scale the sphere term
14 reads -0.74 to -0.80 across samplings that swing the $\|H\|$ partial from -0.29 to
15 $+0.26$; on 86,471 labelled galaxies it is negative on every label, the mismatch be-
16 tween the label’s own Hessian and the shape term is negative wherever the label
17 has local structure, and the shape term alone takes the label-dependent sign that
18 mismatch predicts. Convergence tests must condition on relative curvature, not on
19 one model’s.

20 1 Introduction

21 The Platonic Representation Hypothesis holds that models trained on different data converge toward
22 a shared representation; Duraphe et al. [Duraphe et al., 2025] tested it on the sky: for each of a
23 range of vision models, galaxies imaged by two surveys (Legacy Survey and Hyper Suprime-Cam,
24 HSC) share only 0.3 to 2.3 percent of their nearest neighbours across the two embeddings, a fraction
25 that rises with model capacity, one number per model. We ask where on the manifold that agreement
26 lives: if two models disagree most where the embedding surface bends sharply, local mean curvature
27 should associate negatively with local cross-survey agreement. Testing it needs a pointwise curvature
28 instrument at intrinsic dimension near 20 in ambient dimension 768, and a design that separates
29 curvature from local density.

30 We contribute (i) a pointwise mean-curvature and second-fundamental-form instrument from a
31 trained decoder, validated against analytic truth in the production regime; (ii) a density-controlled
32 test of the curvature–alignment prediction under two probes, which finds no effect free of density;
33 (iii) the residual expansion of a linear map on a submanifold, which names the curvature such a map
34 responds to and shows mean-curvature magnitude blind to it; and (iv) a known-surface and a real-

35 data test on the linear-readout side, where the expansion’s terms behave as predicted and account for
 36 a label-dependent sign.

37 2 Pointwise curvature from a decoder

38 **Data.** DINOv3 ViT-B/16 embeddings of 101,725 galaxies imaged by both HSC and the Legacy
 39 Survey [Duraphe et al., 2025], from which we draw 10,000 rows once with a recorded seed. A sec-
 40 ond set of 86,471 single-survey ViT-B embeddings with catalogue labels (r -band magnitude, pho-
 41 tometric redshift, smooth fraction, stellar mass) serves the readout experiment. Embeddings are
 42 unit-normalized (data on S^{767}); the manifold has trivial first homology ($\beta_1 = 0$ under bootstrap
 43 bands [Fasy et al., 2014]), so one Euclidean chart is appropriate.

44 **Instrument.** Let $F : \mathbb{R}^d \rightarrow \mathbb{R}^{768}$ be the decoder of a plain auto-encoder (three hidden lay-
 45 ers of width 250, SiLU), fit on 80% of rows; curvature is evaluated on held-out rows. With
 46 $J = DF(z)$, metric $g = J^\top J$, normal projector $P_N = I - Jg^{-1}J^\top$ and second fundamen-
 47 tal form $\text{II} = P_N D^2 F(z)$, the mean-curvature vector is $H = \text{tr}_g \text{II}$, by automatic differentiation
 48 checked against finite differences. On the unit sphere every d -submanifold carries a radial compo-
 49 nent of H equal to $-d$ that says nothing about shape, so the map differentiated is $F/\|F\|$ (after
 50 which $H_{\text{rad}} = -d$ to 10^{-14}) and $\|H_{\text{tan}}\|$ carries the shape. Decoder curvature has served as a regular-
 51 izer [Lee and Park, 2023] and been read from templates [Acosta et al., 2022]; here it is an instrument,
 52 validated against closed-form truth.

53 **Validation.** Scored on surfaces with closed-form curvature on four axes: rank (Spearman against
 54 true $\|H\|$), direction (cosine of H vectors), magnitude ratio and calibration. At $D = 768$ an in-sphere
 55 surface $G(z) = Q \text{ normalize}([\phi(z); 0.8 h(z); 0 \dots])$ (ϕ the inverse stereographic map $\mathbb{R}^{16} \rightarrow S^{16}$,
 56 h four Gaussian bumps, Q a fixed rotation, span 18 dimensions) at $n = 86,471$, $k = 2,048$ gives
 57 rank 0.94, direction cosine 0.999, magnitude ratio 1.00, and 0.75 rank under noise at 25% of patch
 58 radius.

59 3 Curvature and cross-survey convergence

60 **Probe.** Per-point mutual k -nearest-neighbour score $\text{MKNN}_i = |N_k(z_1)_i \cap N_k(z_2)_i|/k$ at $k =$
 61 20 between the HSC and Legacy Survey embeddings, the Platonic Universe statistic made local,
 62 correlated with $\|H\|$ over all 10,000 points at $d \in \{20, 25, 32\}$. Centered kernel alignment [Kornblith
 63 et al., 2019] is the second probe.

64 **Result.** The raw association $\rho(\|H\|, \text{MKNN})$ is -0.112 , -0.128 and -0.024 at $d = 20, 25, 32$
 65 ($p < 10^{-5}$, $< 10^{-5}$, 0.009), in the predicted direction. But density correlates with agreement at
 66 -0.21 at every d and with $\|H\|$ at $+0.43$, $+0.32$ and $+0.01$, so residualizing on density leaves
 67 -0.024 , -0.066 and -0.022 : 78%, 49% and 8% of the association removed. Under a permutation
 68 null within density strata only $d = 25$ survives ($p < 5 \times 10^{-5}$; replicating across three decoder seeds;
 69 the CKA tertile gap within strata agrees at $d = 20$ and 25 , -0.05 to -0.08 against a null band of
 70 ± 0.017). Replacing $\|H\|$ by $\|H_{\text{tan}}\|$ collapses the $d = 25$ partial to -0.023 , into the range of the two
 71 d that fail. Convergence is not detectably concentrated by mean-curvature magnitude once density
 72 is controlled.

73 4 What a linear map sees on a curved manifold

74 Let $M \subset \mathbb{R}^D$ be d -dimensional; at x_0 take an orthonormal tangent frame J , the second fundamental
 75 form II (normal-valued, symmetric $d \times d$), $H = \text{tr} \text{II}$, and for a vector w its tangent and normal parts
 76 $w_T + w_N$. Writing M locally as $x_0 + Ju + \frac{1}{2}\text{II}(u, u) + O(|u|^3)$,

$$\nabla_M(w \cdot x) = J^\top w, \quad \text{Hess}_M(w \cdot x) = \langle w_N, \text{II} \rangle, \quad \Delta_M(w \cdot x) = \langle w, H \rangle. \quad (1)$$

77 A linear function is exactly as non-linear on M as M bends in that function’s normal direction. For
 78 $M \subset S^{D-1}$ the normal bundle splits into the sphere’s radial direction and the rest, $\text{II} = \text{II}^S - g \hat{x}$, so
 79 $\text{Hess}_M(w \cdot x) = \langle w_N, \text{II}^S \rangle - (w \cdot \hat{x}) g$ (2): a linear readout of a unit-normalized embedding is bent by




Figure 1: (a) Known surface, intrinsically linear label, exact geometry: three-control partial against local R^2 as the sample’s density–curvature coupling is set by construction (γ from -1 to 1). (b) Galaxies, 512 anchors, $k = 2,048$: partial under the multi-scale density control for the decoder’s $\|H_{\tan}\|$ (filled squares), the shape term (open circles) and the sphere term (open triangles).

the manifold’s shape and, whatever the manifold, by the sphere in proportion to the readout’s own value $\hat{y} - b_0$. For a label y on M , a global affine probe $\hat{y} = w \cdot x + b_0$ and a Gaussian patch with isotropic second moment s^2 in the induced metric, with $\bar{c}$ the patch-mean residual, $\delta = \nabla_M y - J^\top w$, $\Delta = \text{Hess}_M y - \langle w_N, \Pi \rangle$,

$$\mathbb{E}_{\text{patch}}[(y - \hat{y})^2] = \bar{c}^2 + s^2 \|\delta\|^2 + \frac{1}{2} s^4 \|\Delta\|_F^2 + O(s^4 \|\delta\|) + O(s^6). \quad (3)$$

For two embeddings F, G of the same objects and a linear map A aligning their tangent spaces, the same expansion gives $G - (G_0 + A(F - F_0)) = \frac{1}{2}(\Pi_G - A \Pi_F)(u, u) + o(|u|^2)$: the obstruction to linear alignment is the *relative* second fundamental form.

So the curvature a linear map pays for is $\langle w_N, \Pi \rangle$ or $\Pi_G - A \Pi_F$; $\|H\|$ bounds only a trace and need not correlate with either at codimension 750; the term carries s^4 , so patch radius, hence density, multiplies it; and when the label has intrinsic curvature of its own, Δ carries the residual and the curvature piece can take either sign.

5 The probe-facing curvature, on a known surface and on galaxies

The readout side can be tested directly: w is known and Π comes from the decoder. Statistic: five-fold out-of-fold ridge at $\alpha = 100$ from embedding to label; local R^2 inside each of 512 anchors’ 2,048-neighbourhoods; rank-partial Spearman against local R^2 controlling for log patch radius, local label variance and count and, as a stronger control, log radius at $k \in \{16, 64, 256, 1024, 2048\}$. Per anchor we compute $\|H_{\tan}\|$ and the probe-facing Hessian $\|\langle w_N, \Pi \rangle\|_g$ with $w_N = w - Jg^{-1}J^\top w$, split by Eq. (2) into its shape and sphere terms.

Known surface. The validation surface is resampled with weight $\|H_{\tan}\|^\gamma$, so only where the points sit changes (Figure 1a). The $\|H_{\tan}\|$ partial follows the sampling from -0.29 at coupling $+0.80$ through zero to a plateau near $+0.22$, and the decoder reproduces it. The probe-facing Hessian $\|\langle w_N, \Pi \rangle\|$ reads -0.74 to -0.80 at every sampling. On this surface it is the sphere term of Eq. (2): rank 0.996 – 1.000 with $\sqrt{d}|\hat{y} - b_0|$, the in-sphere part 20 to 40 times smaller (one in-sphere normal direction) and, where resolvable, following $\|H_{\tan}\|$. The fixture thus confirms the expansion’s sphere term: an intrinsically linear label loses local R^2 where the readout’s own value is large, whatever the sampling. With a nonlinear label the sphere term reads null and $\|\Delta\|$ carries the partial (-0.2 to -0.6), as Eq. (3) predicts.

Galaxies. Here the two terms of Eq. (2) trade places: the shape term dominates the probe-facing Hessian (rank 0.96 – 0.99 with the full norm). To read it through Eq. (3) we estimate $\text{Hess}_M y$ per anchor by a least-squares quadratic of the label on the tangent-projected coordinates $u = g^{-1}J^\top(x - x_0)$ of its 2,048 neighbours, where the quadratic coefficient is the covariant Hessian. Three regularities (Table 1). The sphere term is negative on every label and both d (-0.13 to -0.36), as on the fixture. Wherever the label has local structure (magnitude, redshift, smooth fraction; stellar mass is null throughout) the mismatch is negative (-0.09 to -0.45) and the alignment $\cos_g(\text{Hess}_M y, \langle w_N, \Pi \rangle)$ positive ($+0.07$ to $+0.43$): local R^2 is higher where the manifold bends in the readout’s direction the way the label bends. The shape term alone is therefore label-signed

label	$\ H_{\text{tan}}\ $		shape $\ \langle w_N, \Pi^S \rangle\ $		sphere $\sqrt{d} \hat{y} - b_0 $		mismatch $\ \Delta\ $		alignment	
	$d=16$	$d=20$	16	20	16	20	16	20	16	20
mag_r	+0.25	-0.17	+0.12	+0.19	-0.26	-0.26	-0.39	-0.40	+0.35	+0.43
photo_z	+0.28	+0.13	-0.06*	-0.12	-0.33	-0.36	-0.45	-0.44	+0.26	+0.24
smooth_fraction	+0.21	+0.15	+0.07*	+0.16	-0.13	-0.15	-0.09*	-0.17	+0.19	+0.07*
stellar_mass	-0.02*	-0.00*	+0.00*	+0.03*	-0.35	-0.35	-0.04*	-0.09*	+0.01*	+0.02*

Table 1: Galaxies, 512 anchors, $k = 2,048$: rank-partial Spearman against local R^2 under the multi-scale density control; alignment is $\cos_g(\text{Hess}_M y, \langle w_N, \Pi \rangle)$. * not significant at 0.05; the three-control partials agree in sign in every significant cell.

(positive for magnitude and smooth fraction, null to weakly negative for redshift) because it is one of three terms in $\|\Delta\|^2 = \|\text{Hess}_M y\|^2 - 2\langle \text{Hess}_M y, \langle w_N, \Pi \rangle \rangle + \|\langle w_N, \Pi \rangle\|^2$: its sign follows the cross term’s where that is nonzero and the remaining square where it is null. $\|H_{\text{tan}}\|$ changes sign between $d = 16$ and 20 for magnitude: it sits on the curve of Figure 1a (coupling to radius -0.60). The same regression applied to the readout $w \cdot x$ returns the point cloud’s own $\langle w_N, \Pi \rangle$ (median metric cosine with the decoder’s 0.55–0.65), and an independent split-half quadratic-chart estimator of Π gives the same shape-term sign on three of four labels.

Across encoders. The quadratic-chart estimator’s mean-bending trace, run on five encoders at the same 512 anchors ($d = 16$, $k = 2,048$: ViT-B, DINOv3, CLIP-B, ConvNeXt-B, ViT-L), gives density-controlled associations with the magnitude probe’s local R^2 of -0.24 , 0.00 , $+0.27$, $+0.23$ and -0.19 . A trace statistic has no encoder-stable sign, as Eq. (1) predicts for a quantity the readout does not see; the terms of Eq. (3) are what a cross-encoder test should carry.

6 Discussion

The design of Section 3 could not have found curvature-concentrated convergence: MKNN and CKA are invariant to linear maps that leave the relative second fundamental form unchanged, so a per-point convergence score can be conditioned on curvature only through $\Pi_G - A \Pi_F$, which needs both embeddings’ second fundamental forms and a map between tangent frames. A pilot at $d = 20$ (2,048 anchors) finds a global ridge map carries HSC tangent frames into the Legacy embedding with 44% residual and no density-free association of $\|\Pi_G(L \cdot, L \cdot) - P_N^G A \Pi_F\|$ with MKNN ($+0.04$, n.s.; a locally fitted tangent map gives -0.18 to -0.24 , a decoder-free estimate $+0.13$): the obstruction is already first order. On the readout side, where the map is known, the sphere term and the Hessian mismatch carry the predicted negative sign on the known surface and on galaxies, and the shape term is label-signed because it enters through its alignment with the label’s own Hessian.

Limitations. One survey pair; the label’s Hessian is a noisy local estimate (R^2 gain over linear 0.10–0.17; a 136-coefficient fit of the same kind failed a split-half reliability gate); the sphere term is not separated from ridge shrinkage of extreme predictions; permutation p -values assume anchor exchangeability under overlapping patches.

References

- Kshitij Duraphe, Michael J. Smith, Shashwat Sourav, and John F. Wu. The platonic universe: Do foundation models see the same sky? In *NeurIPS 2025 Workshop on Machine Learning and the Physical Sciences*, 2025. URL <https://arxiv.org/abs/2509.19453>.
- Brittany Terese Fasy, Fabrizio Lecci, Alessandro Rinaldo, Larry Wasserman, Sivaraman Balakrishnan, and Aarti Singh. Confidence sets for persistence diagrams. *Annals of Statistics*, 42(6):2301–2339, 2014.
- Yonghyeon Lee and Frank Chongwoo Park. On explicit curvature regularization in deep generative models. In *2nd Annual Workshop on Topology, Algebra, and Geometry in Machine Learning (TAG-ML)*, ICML, 2023. URL <https://arxiv.org/abs/2309.10237>.
- Francisco Acosta, Sophia Sanborn, Khanh Dao Duc, Manu Madhav, and Nina Miolane. Quantifying extrinsic curvature in neural manifolds. *arXiv preprint*, 2022. URL <https://arxiv.org/abs/2212.10414>.

154 Simon Kornblith, Mohammad Norouzi, Honglak Lee, and Geoffrey Hinton. Similarity of neural network rep-
155 resentations revisited. In *Proceedings of the 36th International Conference on Machine Learning (ICML)*,
156 2019.